

Exponential Family

- Suppose Y_1, \dots, Y_n are independent random variables.
- Let $f(y_i; \theta_i, \phi)$ be the Probability Mass Function (PMF) or Probability Density Function (PDF) of Y_i , where ϕ is a scale parameter.

- If we can write

$$f(y_i; \theta_i, \phi) = \exp \left(y_i \theta_i - b(\theta_i) \frac{1}{a(\phi)} + c(y_i, \phi) \right),$$

then we call the PMF or the PDF $f(y_i; \theta_i, \phi)$ an exponential family.

Problem 1.

Find the form of GLM for the following distributions, and show the reasonable link function:

1. Normal distributions
2. Inverse Gaussian
3. Binomial distribution
4. Poisson distribution
5. Gamma distribution
6. Beta

Normal Distribution

Assume $Y_i \sim \mathcal{N}(\mu_i, \sigma^2)$. Then, $E(Y_i) = \mu_i$ and σ is a scale parameter. The Probability Density Function (PDF) is given by

$$\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y_i - \mu_i)^2}{2\sigma^2}} = \exp \left\{ \frac{y_i \mu_i - \mu_i^2/2}{\sigma^2} - \left(\frac{1}{2} \log(2\pi\sigma^2) - \frac{y_i^2}{2\sigma^2} \right) \right\}.$$

Thus, use

- $\theta_i = \mu_i$,
- $b(\theta_i) = \frac{\theta_i^2}{2}$,
- $\phi = \sigma^2$,
- $a(\phi) = \phi$,
- $c(y_i, \phi) = -\frac{1}{2} \log(2\pi\phi) - \frac{y_i^2}{2\phi}$.

$$E[Y_{ij} | x_{ij}] = h(x_{ij}^T \beta) = x_{ij}^T \beta$$

Inverse Gaussian Distribution

Let us rewrite the probability density function (pdf) of the Inverse Gaussian distribution with parameters μ_i and λ :

$$f(y_i; \mu_i, \lambda) = \left(\frac{\lambda}{2\pi y_i^3} \right)^{1/2} \exp \left\{ -\frac{\lambda(y_i - \mu_i)^2}{2\mu_i^2 y_i} \right\}, \quad y > 0$$

in the following form:

$$\begin{aligned} f(y_i; \mu_i, \lambda) &= \exp \left\{ -\frac{\lambda(y_i - \mu_i)^2}{2\mu_i^2 y_i} + \frac{1}{2} \ln \left(\frac{\lambda}{2\pi y_i^3} \right) \right\} \\ &= \exp \left\{ \frac{-\frac{1}{2\mu_i^2} y_i + \frac{1}{\mu_i}}{\frac{1}{\lambda}} + \left(\frac{1}{2} \ln \frac{\lambda}{2\pi y_i^3} - \frac{\lambda}{2y_i} \right) \right\} \end{aligned}$$

Now, let's identify the exponential family components:

- Canonical parameter: $\theta_i = -\frac{1}{2\mu_i^2}$
- $b(\theta_i) = -\frac{1}{\mu_i} = -(-2\theta_i)^{\frac{1}{2}}$
- $\phi = \lambda$
- $a(\phi) = \frac{1}{\phi}$
- $c(y_i, \phi) = \frac{1}{2} \ln \frac{\phi}{2\pi y_i^3} - \frac{\phi}{2y_i}$

Thus, the Inverse Gaussian distribution can be shown to be a member of the exponential family.

$$E[Y_{ij} | x_{ij}] = h(x_{ij}^T \beta) = \exp(x_{ij}^T \beta) > 0$$

Binomial Distribution

Assume $Y_i \sim \text{Bin}(n_i, p_i)$. Then, $E(Y_i) = n_i p_i$. The Probability Mass Function (PMF) is given by

$$\binom{n_i}{y_i} p_i^{y_i} (1 - p_i)^{n_i - y_i} = \exp \left\{ y_i \log \left(\frac{p_i}{1 - p_i} \right) + n_i \log(1 - p_i) - \log \binom{n_i}{y_i} \right\}.$$

Thus,

- $\theta_i = \log \left(\frac{p_i}{1 - p_i} \right)$,
- $b(\theta_i) = n_i \log(1 + e^{\theta_i})$,
- $\phi = 1$, $a(\phi) = 1$,
- $c(y, \phi) = -\log \binom{n_i}{y_i}$.

$$E[Y_{ij} | x_{ij}] = h(x_{ij}^T \beta) = n \frac{\exp(x_{ij}^T \beta)}{1 + \exp(x_{ij}^T \beta)} \in (0, n)$$

Poisson Distribution

Assume $Y_i \sim \text{Poisson}(\lambda_i)$. Then, $E(Y_i) = \lambda_i$. The Probability Mass Function (PMF) is given by

$$\frac{\lambda_i^{y_i} e^{-\lambda_i}}{y_i!} = \exp \{y_i \log(\lambda_i) - \lambda_i - \log(y_i!)\}.$$

Thus,

- $\theta_i = \log(\lambda_i)$,
- $b(\theta_i) = e^{\theta_i}$,
- $\phi = 1$, $a(\phi) = 1$,
- $c(y_i, \phi) = -\log(y_i!)$.

$$E[Y_{ij} | x_{ij}] = h(x_{ij}^T \beta) = \exp(x_{ij}^T \beta) > 0$$

Gamma Distribution

Assume $x_i \sim \Gamma(\alpha, \beta_i)$, where β_i is unknown. Then, $E(x_i) = \frac{\alpha}{\beta_i}$. The Probability Mass Function (PMF) is given by

$$\frac{\beta_i^\alpha x_i^{\alpha-1} e^{-\beta_i x_i}}{\Gamma(\alpha)} = \exp \{ \alpha \log x_i + \alpha \log(\beta_i) - \log(\Gamma(\alpha)) - \log(x_i) - \beta_i x_i \}.$$

Assuming α is known, if we choose $y_i = x_i$, then

- $\theta_i = -\beta_i$ ($\theta_i < 0$),
- $b(\theta_i) = -\alpha \log(-\theta_i)$,
- $\phi = 1$, and $a(\phi) = 1$.

Remark: We can also choose $y_i = -x_i$ and $\theta_i = \beta_i$. In this case, $b(\theta_i) = -\alpha \log(\theta_i)$.

$$E[Y_{ij} | x_{ij}] = h(x_{ij}^T \beta) = \exp(x_{ij}^T \beta) > 0$$

Problem 2.

Paper Summarization

The part 1 of this page talks about:

- **Sliced inverse regression (SIR)**: A novel data-analytic tool for reducing the dimension of the input variable x without fitting any parametric or nonparametric model¹[1]. It explores the inverse view of regression, where x is regressed against y , and uses a simple step function to estimate the inverse regression curve²[2]. - **Effective dimension reduction (e.d.r.) space**: The linear space generated by the unknown row vectors $3k$ ($k = 1, \dots, K$) in the model $y = f(3lx, \dots, 3Kx, e)$, where f is an arbitrary unknown function. The goal is to estimate this space, which captures all the information about y from x . - **Inverse regression curve**: The curve $E(x | y)$ that connects the conditional mean of x given y as y varies. Under certain conditions, this curve falls into the e.d.r. space. A principal component analysis on the covariance matrix of the estimated inverse regression curve can locate its main orientation, yielding the estimates for e.d.r. directions³[3]. - **Sampling properties of SIR**: The output of SIR provides root n consistent estimates for the e.d.r. directions under a design condition on the distribution of x [4]. The eigenvalues of the covariance matrix can be used to assess the number of components in the model and the effectiveness of SIR. - **Simulation results**: SIR is demonstrated to be effective in reducing the dimension of x from 10 to 2 for a data set with 400 observations. The spin-plot of y against the projected variables obtained by SIR mimics the spin-plot of y against the true directions very well [5]. A chi-squared statistic is proposed to test whether a direction found by SIR is spurious [6].

Problem 3.

Binary response ($Y_{ij} = 0/1$)

Logistic regression model:

1. $\mathbb{E}[Y_{ij} | x_{ij}] = h(x_{ij}^T \beta) = \frac{\exp(x_{ij}^T \beta)}{1 + \exp(x_{ij}^T \beta)}$
2. $\mathbf{Var}[Y_{ij} | x_{ij}] = \mathbb{E}[Y_{ij} | x_{ij}](1 - \mathbb{E}[Y_{ij} | x_{ij}])$
3. $\mathbf{Cor}[Y_{ij}, Y_{ik} | x_{ij}, x_{ik}] = \alpha$
4. odd ratio: $OR(Y_{ij}, Y_{ik}) = \frac{P(Y_{ij}=1, Y_{ik}=1)P(Y_{ij}=0, Y_{ik}=0)}{P(Y_{ij}=1, Y_{ik}=0)P(Y_{ij}=0, Y_{ik}=1)}$

$$\begin{aligned}
 \alpha = \mathbf{Cor}[Y_{ij}, Y_{ik} | x_{ij}, x_{ik}] &= \frac{\mathbb{E}[(Y_{ij} - \mathbb{E}[Y_{ij} | x_{ij}])(Y_{ik} - \mathbb{E}[Y_{ik} | x_{ik}])]}{\sigma_{Y_{ij} | x_{ij}} \sigma_{Y_{ik} | x_{ik}}} \\
 &= \frac{\mathbb{E}[(Y_{ij} - \mathbb{E}[Y_{ij} | x_{ij}])(Y_{ik} - \mathbb{E}[Y_{ik} | x_{ik}])]}{\sqrt{\mathbf{Var}[Y_{ij} | x_{ij}]} \sqrt{\mathbf{Var}[Y_{ik} | x_{ik}]}} \\
 &= \frac{\mathbb{E}[(Y_{ij} - \mathbb{E}[Y_{ij} | x_{ij}])(Y_{ik} - \mathbb{E}[Y_{ik} | x_{ik}])]}{\sqrt{\mathbb{E}[Y_{ij} | x_{ij}](1 - \mathbb{E}[Y_{ij} | x_{ij}])} \sqrt{\mathbb{E}[Y_{ik} | x_{ik}](1 - \mathbb{E}[Y_{ik} | x_{ik}])}} \\
 &= \frac{\mathbb{E}[Y_{ij}Y_{ik} - Y_{ij}h(x_{ik}^T \beta) - Y_{ik}h(x_{ij}^T \beta) + h(x_{ij}^T \beta)h(x_{ik}^T \beta)]}{\sqrt{h(x_{ij}^T \beta)(1 - h(x_{ij}^T \beta))} \sqrt{h(x_{ik}^T \beta)(1 - h(x_{ik}^T \beta))}} \\
 &= \frac{\mathbb{E}[Y_{ij}Y_{ik}] - h(x_{ij}^T \beta)h(x_{ik}^T \beta)}{\sqrt{h(x_{ij}^T \beta)(1 - h(x_{ij}^T \beta))} \sqrt{h(x_{ik}^T \beta)(1 - h(x_{ik}^T \beta))}} \\
 &= \frac{P(Y_{ij} = 1, Y_{ik} = 1) - h(x_{ij}^T \beta)h(x_{ik}^T \beta)}{\sqrt{h(x_{ij}^T \beta)(1 - h(x_{ij}^T \beta))} \sqrt{h(x_{ik}^T \beta)(1 - h(x_{ik}^T \beta))}} \in [0, 1]
 \end{aligned}$$

.....

$$\begin{aligned}
 \gamma = OR(Y_{ij}, Y_{ik}) &= \frac{P(Y_{ij} = 1, Y_{ik} = 1)P(Y_{ij} = 0, Y_{ik} = 0)}{P(Y_{ij} = 1, Y_{ik} = 0)P(Y_{ij} = 0, Y_{ik} = 1)} \\
 &= \frac{P_{ij,ik}[1 - (h(x_{ij}^T \beta) - P_{ij,ik}) - (h(x_{ik}^T \beta) - P_{ij,ik}) - P_{ij,ik}]}{[h(x_{ij}^T \beta) - P_{ij,ik}][h(x_{ik}^T \beta) - P_{ij,ik}]}
 \end{aligned}$$

Solve for $P_{ij,ik}$ using γ and $h(\cdot)$.

Problem 4.

Describe

1. how to conduct the EM (Expectation-Maximization) algorithm
2. how to conduct MCMC

Denoting

$$Q(q|q_0, x) = \mathbb{E}_{q_0} [\log L_c(q|x, Z)],$$

the EM algorithm indeed proceeds "iteratively" by maximizing $Q(q|q_0, x)$ at each iteration, and, if $q^{(\hat{1})}$ is the value of q maximizing $Q(q|q_0, x)$, by replacing q_0 by the updated value $q^{(\hat{1})}$. In this manner, a sequence of estimators $\{q^{(\hat{j})}\}_j$ is obtained, where

$$Q(q^{(\hat{j})}|q^{(\hat{j-1})})$$

Pick a starting value $q^{(\hat{0})}$ and set $m = 0$. Repeat

1. Compute (the E-step)

$$Q(q|q^{(\hat{m}),x}) = \mathbb{E}_{q^{(\hat{m})}} [\log L_c(q|x, Z)],$$

where the expectation is with respect to $k(z|q^{(\hat{m}),x})$.

2. Maximize $Q(q|q^{(\hat{m}),x})$ in q and take (the M-step)

$$q^{(\hat{m+1})} = \arg \max_q Q(q|q^{(\hat{m}),x})$$

and set $m = m + 1$.

until a fixed point is reached; i.e., $q^{(\hat{m+1})} = q^{(\hat{m})}$.

.....

Gibbs Sampling of MCMC (Markov Chain Monte Carlo)

Consider $f(y_{ij} | u_i, \beta) = e^{\left(\frac{y_{ij}\theta_{ij} - \psi(\theta_{ij})}{\alpha(\phi)} - c(y_{ij}; \phi)\right)}$

with $g(u_i | G) = (2\pi)^{-\frac{q}{2}} |G|^{-\frac{1}{2}} e^{-\frac{u_i^T G^{-1} u_i}{2}}$, and $h(\mu_{ij}) = x_{ij}^T \beta + z_{ij}^T u_i$.

The likelihood function of (β, G) is

$$L(\beta, G | y) \propto \prod_{i=1}^n \prod_{j=1}^{m_i} f(y_{ij} | u_i, \beta) |G|^{-1/2} \exp\left\{-\frac{u_i^T G^{-1} u_i}{2}\right\} du_i.$$

In a Bayesian approach to analyzing the random effects GLM, the parameters (β, G) are random variables and are treated symmetrically with the longitudinal measurements and unobserved latent variables. Thus, the random effects GLM is an example of a hierarchical Bayes model.

Assumptions: $[\beta | G, U, y] = [\beta | U, y]$, $[G | \beta, U, y] = [G | U]$ and $[U | \beta, G, y]$.

1. Assume that β has a flat prior function. Then,

$[\beta | U^{(k)}, y] \propto \prod_{i=1}^n \prod_{j=1}^{m_i} f(y_{ij} | U_i^{(k)}, \beta) \approx N(\beta^{(k)}, V_\beta^{(k)})$, as $n \rightarrow \infty$, where $\beta^{(k)}$ is the maximum likelihood estimator and $V_\beta^{(k)}$ is the inverse of the Fisher information.

Adjustment for smaller samples - "Rejection sampling" (Ripley, 1987)

Let $f\left(\beta \mid U^{(k)}, y\right)$ and $\phi\left(\beta \mid \beta^{(k)}, V_{\beta}^{(k)}\right)$ denote separately the true density and Gaussian density.

Choose a constant $c \geq 1$ such that $c\phi\left(\beta \mid \beta^{(k)}, V_{\beta}^{(k)}\right) \geq f\left(\beta \mid U^{(k)}, y\right)$.

Step1: Generate $\beta^* \sim \phi\left(\beta \mid \beta^{(k)}, V_{\beta}^{(k)}\right)$ and $u \sim U(0, 1)$.

Step2: If $\frac{f\left(\beta^* \mid U^{(k)}, y\right)}{c\phi\left(\beta^* \mid \beta^{(k)}, V_{\beta}^{(k)}\right)} < u$, $\beta^{(k+1)} = \beta^*$. Otherwise, the process returns to Step1.

2. $[G \mid U^{(k)}]$

Assume that $\pi(G) \propto |G|^{-1}$: non-informative prior (see Box and Tiao, 1973).

Then, $[G \mid U^{(k)}] \sim$ Inverted Wishart $(S^{(k)}, n - q + 1)$, where $S^{(k)} = \sum_{i=1}^n U_i^{(k)} U_i^{(k)T}$

Remark

If $A \sim$ Wishart $(\Sigma_{p \times p}, n)$, the p.d.f of A is $f_A(A) \propto |A|^{\frac{-1}{2}(n-p-1)} e^{\frac{-1}{2}tr \Sigma^{-1}A}$.

It implies that $B = A^{-1} \sim$ Inverted Wishart (Σ^{-1}, n) with p.d.f. $f_B(B) \propto |B|^{\frac{-1}{2}(n+p+1)} e^{\frac{-1}{2}tr \Sigma^{-1}B^{-1}}$.

Thus,

$$\pi\left(G \mid U^{(k)}\right) \propto |G|^{\frac{-1}{2}(n+2)} e^{\frac{-1}{2}tr(S^{(k)}G^{-1})},$$

i.e. $[G \mid U^{(k)}] \sim$ Inverted Wishart $(S^{(k)}, n - q + 1)$.

3. $[U \mid \beta^{(k)}, G^{(k)}, y]$

Using

$$f\left(U_i \mid \hat{\beta}^{(k)}, G^{(k)}, y_i\right) \propto f\left(y_i \mid U_i, \hat{\beta}^{(k)}\right) g\left(U_i \mid G^{(k)}\right) \triangleq f_n\left(U_i \mid \hat{\beta}^{(k)}, G^{(k)}, y_i\right),$$

we can find the mode and curvature of $f_n\left(U_i \mid \hat{\beta}^{(k)}, G^{(k)}, y_i\right)$, which matches a Gaussian density.

Using the surrogate response

$$Z_i^* = X_i \beta + D_i U_i + \text{Diag}\left(h'(\mu_i)\right)(y_i - \mu_i),$$

the maximum value of $f_n\left(U_i \mid \hat{\beta}^{(k)}, G^{(k)}, y_i\right)$ occurs at

$$U_i = \left(D_i^T Q_i^{-1} D_i + G^{(k)-1}\right)^{-1} D_i^T Q_i^{-1} \left(Z_i^* - X_i \beta^{(k)}\right) = G^{(k)} D_i \left(D_i G^{(k)} D_i^T + Q_i\right)^{-1} \left(Z_i^* - X_i \hat{\beta}^{(k)}\right)$$

and its curvature is $V_i = \left(D_i^T Q_i^{-1} D_i + G^{(k)-1}\right)^{-1}$. Similar to the method in (3), $U_i^{(k)}$ can be obtained.

Problem 5.

Theorem 3.1.

Under the regularity conditions, $n^{\frac{-1}{2}} \left[\begin{pmatrix} \hat{\beta} \\ \hat{\alpha} \end{pmatrix} - \begin{pmatrix} \beta \\ \alpha \end{pmatrix} \right] \xrightarrow{d} N(0, \Sigma)$, where Σ can be estimated by $(\frac{1}{n} \sum_{i=1}^n C_i^T B_i^{-1} D_i)^{-1} (\frac{1}{n} \sum_{i=1}^n C_i^T B_i^{-1} V_{0i} D_i B_i^{-1} C_i) (\frac{1}{n} \sum_{i=1}^n D_i^T B_i^{-1} C_i)^{-1}$, where $C_i = \begin{pmatrix} \frac{\partial \mu_i}{\partial \beta} & 0 \\ 0 & \frac{\partial \eta_i}{\partial \alpha} \end{pmatrix}$, $B_i = \begin{pmatrix} \text{Var}(Y_i) & 0 \\ 0 & H_i \end{pmatrix}$, $D_i = \begin{pmatrix} \frac{\partial \mu_i}{\partial \beta} & \frac{\partial \mu_i}{\partial \alpha} \\ \frac{\partial \eta_i}{\partial \beta} & \frac{\partial \eta_i}{\partial \alpha} \end{pmatrix}$, and $V_{0i} = \begin{pmatrix} y_i - \mu_i \\ \omega_i - \eta_i \end{pmatrix}^{\otimes 2}$.